

In [6]: `import torch`

```
for i in range(torch.cuda.device_count()):
    print(f"CUDA:{i} -> {torch.cuda.get_device_name(i)}")
```

CUDA:0 -> NVIDIA GeForce RTX 3050 Laptop GPU

In [1]: `import torch`
`import torch.nn as nn`
`import torch.optim as optim`
`import torchvision`
`import torchvision.transforms as transforms`
`import torchvision.models as models`
`import time`
`import matplotlib.pyplot as plt`
`import numpy as np`

`torch.backends.cudnn.benchmark = True`

`device = torch.device("cuda:0")`
`print(f"Using device: {device}")`

`transform_train = transforms.Compose([`
 `transforms.Resize(224),`
 `transforms.RandomHorizontalFlip(),`
 `transforms.RandomRotation(10),`
 `transforms.ToTensor(),`
 `transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])`
`])`

`transform_test = transforms.Compose([`
 `transforms.Resize(224),`
 `transforms.ToTensor(),`
 `transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])`
`])`

`batch_size = 32`
`subset_ratio = 0.1`
`dataset_path = "./data"`

`trainset = torchvision.datasets.CIFAR10(root=dataset_path, train=True, download=False, transform=transform_train)`
`trainset = torch.utils.data.Subset(trainset, range(int(len(trainset) * subset_ratio)))`
`trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, shuffle=True, num_workers=4)`

`testset = torchvision.datasets.CIFAR10(root=dataset_path, train=False, download=False, transform=transform_test)`
`testset = torch.utils.data.Subset(testset, range(int(len(testset) * subset_ratio)))`
`testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=False, num_workers=4)`

`classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')`

`vgg16 = models.vgg16(weights=models.VGG16_Weights.DEFAULT)`
`vgg16.classifier[6] = nn.Linear(vgg16.classifier[6].in_features, 10)`
`vgg16 = vgg16.to(device)`

`criterion = nn.CrossEntropyLoss()`
`optimizer = optim.SGD(vgg16.parameters(), lr=0.001, momentum=0.9)`
`scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'min', patience=2, factor=0.5)`

`scaler = torch.cuda.amp.GradScaler()`

```

def train_model(model, epochs=5):
    train_accuracies, val_accuracies, val_losses = [], [], []

    for epoch in range(epochs):
        start_time = time.time()
        model.train()
        running_loss, correct, total = 0.0, 0, 0

        for i, data in enumerate(trainloader, 0):
            inputs, labels = data[0].to(device, non_blocking=True), data[1].to(device, non_blocking=True)

            optimizer.zero_grad()
            with torch.cuda.amp.autocast():
                outputs = model(inputs)
                loss = criterion(outputs, labels)
            scaler.scale(loss).backward()
            scaler.step(optimizer)
            scaler.update()

            running_loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

            if i % 20 == 19:
                print(f'[{epoch + 1}, {i + 1}] loss: {running_loss / 20:.3f}')
                running_loss = 0.0

        train_accuracy = 100 * correct / total
        train_accuracies.append(train_accuracy)

        model.eval()
        val_loss, correct, total = 0.0, 0, 0

        with torch.no_grad():
            for data in testloader:
                images, labels = data[0].to(device, non_blocking=True), data[1].to(device, non_blocking=True)
                with torch.cuda.amp.autocast():
                    outputs = model(images)
                    loss = criterion(outputs, labels)
                val_loss += loss.item()
                _, predicted = torch.max(outputs.data, 1)
                total += labels.size(0)
                correct += (predicted == labels).sum().item()

        val_accuracy = 100 * correct / total
        val_accuracies.append(val_accuracy)
        avg_val_loss = val_loss / len(testloader)
        val_losses.append(avg_val_loss)
        scheduler.step(avg_val_loss)

        print(f'Epoch {epoch + 1}/{epochs} completed in {time.time() - start_time:.1f}s')
        print(f'Train Accuracy: {train_accuracy:.2f}%, Val Accuracy: {val_accuracy:.2f}%')
        print(f'Learning Rate: {optimizer.param_groups[0]["lr"]:.6f}')
        print('-' * 60)

    return train_accuracies, val_accuracies, val_losses

print("Starting training...")
epochs = 5
train_accuracies, val_accuracies, val_losses = train_model(vgg16, epochs)

```

```
torch.save(vgg16.state_dict(), 'vgg16_cifar10.pth')  
print("Training complete. Model saved.")
```

Using device: cuda:0

```
C:\Users\agnik\AppData\Local\Temp\ipykernel_28932\4036943035.py:52: FutureWarning: `torch.cuda.amp.  
GradScaler(args...)` is deprecated. Please use `torch.amp.GradScaler('cuda', args...)` instead.  
    scaler = torch.cuda.amp.GradScaler()
```

Starting training...

```
C:\Users\agnik\AppData\Local\Temp\ipykernel_28932\4036943035.py:66: FutureWarning: `torch.cuda.amp.  
autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.  
    with torch.cuda.amp.autocast():
```

```
[1, 20] loss: 2.165  
[1, 40] loss: 1.429  
[1, 60] loss: 1.025  
[1, 80] loss: 0.888  
[1, 100] loss: 0.792  
[1, 120] loss: 0.692  
[1, 140] loss: 0.683
```

```
C:\Users\agnik\AppData\Local\Temp\ipykernel_28932\4036943035.py:91: FutureWarning: `torch.cuda.amp.  
autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.  
    with torch.cuda.amp.autocast():
```

Epoch 1/5 completed in 204.2s
Train Accuracy: 62.64%, Val Accuracy: 80.40%
Learning Rate: 0.001000

[2, 20] loss: 0.512
[2, 40] loss: 0.555
[2, 60] loss: 0.490
[2, 80] loss: 0.524
[2, 100] loss: 0.597
[2, 120] loss: 0.472
[2, 140] loss: 0.376
Epoch 2/5 completed in 437.1s
Train Accuracy: 82.80%, Val Accuracy: 83.40%
Learning Rate: 0.001000

[3, 20] loss: 0.426
[3, 40] loss: 0.284
[3, 60] loss: 0.425
[3, 80] loss: 0.422
[3, 100] loss: 0.353
[3, 120] loss: 0.350
[3, 140] loss: 0.406
Epoch 3/5 completed in 431.0s
Train Accuracy: 86.90%, Val Accuracy: 85.50%
Learning Rate: 0.001000

[4, 20] loss: 0.304
[4, 40] loss: 0.358
[4, 60] loss: 0.288
[4, 80] loss: 0.271
[4, 100] loss: 0.266
[4, 120] loss: 0.244
[4, 140] loss: 0.281
Epoch 4/5 completed in 429.6s
Train Accuracy: 89.80%, Val Accuracy: 84.30%
Learning Rate: 0.001000

[5, 20] loss: 0.327
[5, 40] loss: 0.229
[5, 60] loss: 0.189
[5, 80] loss: 0.225
[5, 100] loss: 0.212
[5, 120] loss: 0.245
[5, 140] loss: 0.259
Epoch 5/5 completed in 431.9s
Train Accuracy: 91.70%, Val Accuracy: 85.50%
Learning Rate: 0.001000

Training complete. Model saved.

```
In [2]: plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(range(1, epochs+1), train_accuracies, label='Training Accuracy')
plt.plot(range(1, epochs+1), val_accuracies, label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy (%)')
plt.title('Model Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(range(1, epochs+1), val_losses)
plt.xlabel('Epochs')
```

```

plt.ylabel('Loss')
plt.title('Validation Loss')

plt.tight_layout()
plt.savefig('vgg16_performance.png')

def evaluate_model(model, dataloader):
    model.eval()
    correct, total = 0, 0
    class_correct = list(0. for _ in range(10))
    class_total = list(0. for _ in range(10))

    with torch.no_grad():
        for data in dataloader:
            images, labels = data[0].to(device, non_blocking=True), data[1].to(device, non_blocking=True)
            outputs = model(images)
            _, predicted = torch.max(outputs, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

            c = (predicted == labels).squeeze()
            for i in range(labels.size(0)):
                class_correct[labels[i]] += c[i].item()
                class_total[labels[i]] += 1

    print(f'Overall Accuracy on test images: {100 * correct / total:.2f}%')
    for i in range(10):
        print(f'Accuracy of {classes[i]}: {100 * class_correct[i] / class_total[i]:.2f}%')

    return 100 * correct / total

print("\nEvaluating model on test data...")
test_accuracy = evaluate_model(vgg16, testloader)

def visualize_predictions(model, dataloader, classes, num_images=5):
    model.eval()
    dataiter = iter(dataloader)
    images, labels = next(dataiter)

    with torch.no_grad():
        outputs = model(images[:num_images].to(device))
        _, predicted = torch.max(outputs, 1)
        predicted = predicted.cpu().numpy()

    images = images[:num_images]
    labels = labels[:num_images]

    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])

    fig, axes = plt.subplots(1, num_images, figsize=(15, 3))
    for i in range(num_images):
        img = images[i].numpy().transpose((1, 2, 0))
        img = std * img + mean
        img = np.clip(img, 0, 1)

        axes[i].imshow(img)
        axes[i].set_title(f"True: {classes[labels[i]]}\nPred: {classes[predicted[i]]}")
        axes[i].axis('off')

    plt.tight_layout()
    plt.savefig('predictions.png')

```

```
print("\nVisualizing some predictions...")
visualize_predictions(vgg16, testloader, classes)
```

Evaluating model on test data...

Overall Accuracy on test images: 85.50%

Accuracy of plane: 84.47%

Accuracy of car: 95.51%

Accuracy of bird: 90.00%

Accuracy of cat: 73.79%

Accuracy of deer: 76.67%

Accuracy of dog: 83.72%

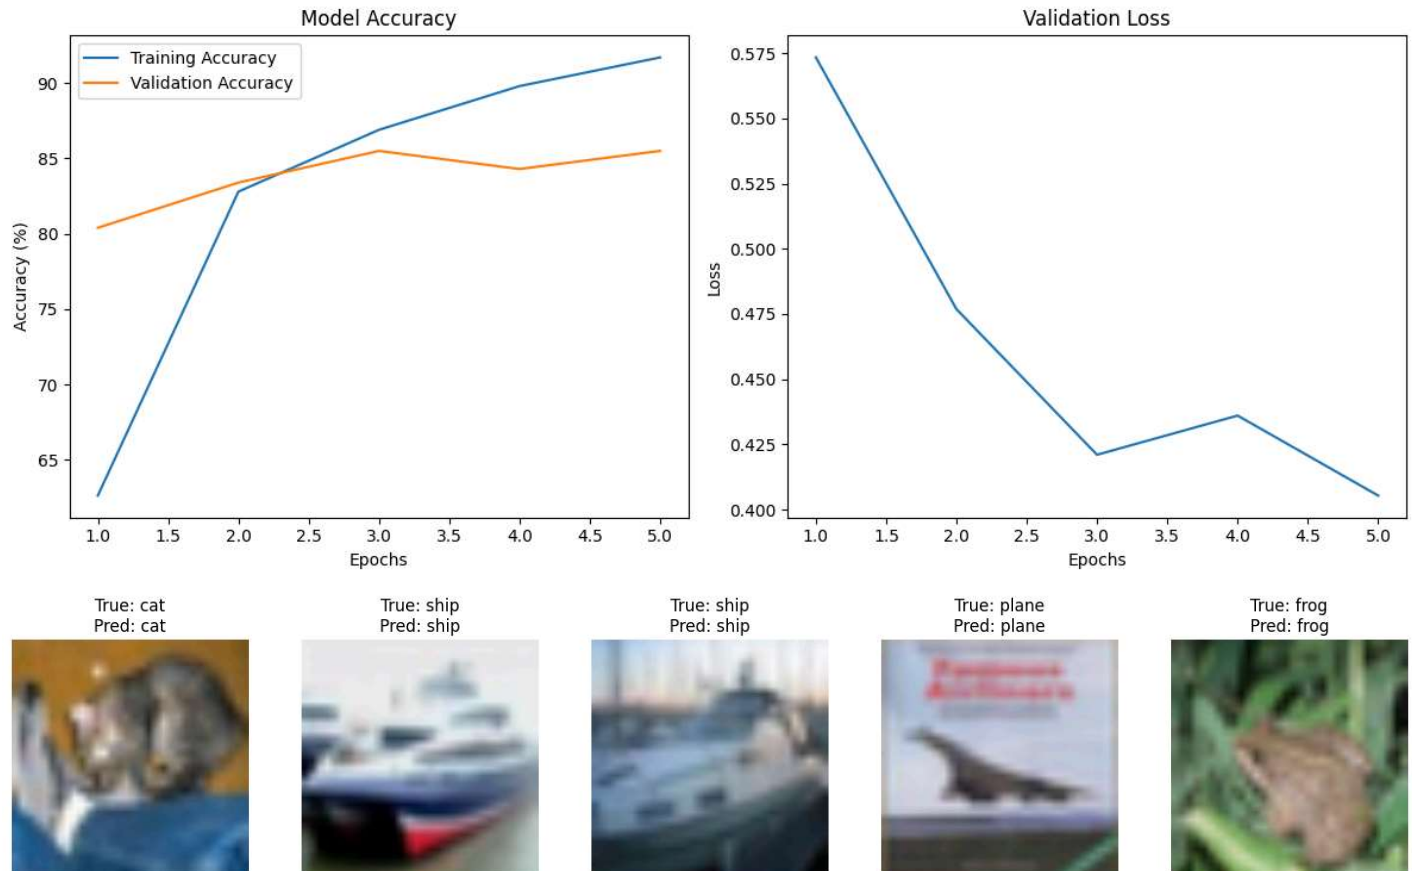
Accuracy of frog: 90.18%

Accuracy of horse: 85.29%

Accuracy of ship: 88.68%

Accuracy of truck: 86.24%

Visualizing some predictions...



```
In [5]: from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns

def evaluate_with_metrics(model, dataloader, classes):
    model.eval()
    all_preds = []
    all_labels = []

    with torch.no_grad():
        for data in dataloader:
            images, labels = data[0].to(device, non_blocking=True), data[1].to(device, non_blocking=True)
            outputs = model(images)
            _, predicted = torch.max(outputs, 1)

            all_preds.extend(predicted.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())
```

```

print("\nClassification Report:\n")
print(classification_report(all_labels, all_preds, target_names=classes, digits=4))

cm = confusion_matrix(all_labels, all_preds)
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=classes, yticklabels=classes)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.tight_layout()
plt.savefig('confusion_matrix.png')
plt.show()


print("\nGenerating evaluation metrics...")
evaluate_with_metrics(vgg16, testloader, classes)

```

Generating evaluation metrics...

Classification Report:

	precision	recall	f1-score	support
plane	0.8529	0.8447	0.8488	103
car	0.8252	0.9551	0.8854	89
bird	0.7826	0.9000	0.8372	100
cat	0.8085	0.7379	0.7716	103
deer	0.9200	0.7667	0.8364	90
dog	0.7423	0.8372	0.7869	86
frog	0.8559	0.9018	0.8783	112
horse	0.9062	0.8529	0.8788	102
ship	0.9307	0.8868	0.9082	106
truck	0.9495	0.8624	0.9038	109
accuracy			0.8550	1000
macro avg	0.8574	0.8545	0.8535	1000
weighted avg	0.8599	0.8550	0.8552	1000



In []: